

---

# Beyond Graph Recovery: Implicit Causal Discovery at Scale

---

Anonymous Authors<sup>1</sup>

## Abstract

Traditional causal discovery methods require recovering the full causal graph before predicting intervention effects, a two-stage approach that struggles to scale beyond  $\sim 20$  variables. We introduce an implicit causal discovery method using a transformer architecture inspired by Prior-Data Fitted Networks (Müller et al., 2021). By employing TabPFN-style embeddings (Hollmann et al., 2022)—specifically, an MLP-based value encoder and interleaved [Feature, Value] token strategy—our model learns to predict intervention outcomes in a single forward pass, bypassing explicit graph recovery and combinatorial search. On interventional prediction tasks, our approach is comparable to the PC Algorithm at 20 variables (1.34 vs 1.37 MAE) and significantly outperforms it as systems scale beyond 30 variables, offering superior RMSE (4.71 vs 4.91) and real-time inference (3.9ms). Through systematic evaluation on 5,000 zero-shot test cases across unseen graph structures, we demonstrate successful scaling to 50 variables—a regime where traditional methods often fail. A scaling analysis reveals that capacity limits stem from  $O(N^2)$  attention complexity rather than parameter count, suggesting sparse attention as a promising path for further scaling. Our work demonstrates that for prediction-focused applications, implicit learning provides a practical, scalable alternative to explicit graph recovery.

## 1. Introduction

Consider a healthcare system monitoring 40 patient biomarkers. A clinician wants to predict how intervening on blood pressure medication will affect other vital signs. Traditional causal discovery would first attempt to learn the

<sup>1</sup>Anonymous Institution, Anonymous City, Anonymous Region, Anonymous Country. Correspondence to: Anonymous Author <anon.email@domain.com>.

Preliminary work. Under review by the International Conference on Machine Learning (ICML). Do not distribute.

network structure among all  $\binom{40}{2} = 780$  potential pairwise interactions—a search space involving  $2^{780}$  possible directed acyclic graph (DAG) structures. We ask: is this explicit graph recovery step necessary if our primary goal is prediction?

### 1.1. The Computational Burden and the "Two-Stage" Trap

Causal discovery traditionally follows a rigid pipeline: (1) learn the causal graph  $\mathcal{G}$ , then (2) use  $\mathcal{G}$  to estimate intervention effects. This paradigm faces three fundamental challenges:

**Computational complexity.** Graph search is NP-hard (Chickering, 1996). Even "soft" optimization methods like NOTEARS (Zheng et al., 2018) struggle with the cubic scaling of gradient-based DAG learning beyond 30 variables.

**Identifiability and Error Propagation.** Interventional data often cannot uniquely identify  $\mathcal{G}$ , yet intervention effects may still be computable (Pearl, 2009). Furthermore, small errors in graph recovery can lead to significant bias in effect estimation.

**The Scalability Gap.** High-stakes applications—from personalized medicine to economic policy—require predictions across 40–100 variables, a regime where traditional recovery-based methods are computationally intractable.

### 1.2. Implicit Causal Discovery

We propose a fundamental departure: *skip graph recovery entirely* and learn the intervention operator directly. Our key insight is that transformer attention mechanisms can implicitly capture causal dependencies without explicit DAG constraints.

**Core idea.** Instead of learning  $\mathcal{G}$  then using it to compute  $\mathbb{E}[X \mid \text{do}(X_i = v)]$ , we directly learn:

$$f : (\mathbf{X}_{\text{obs}}, \text{intervention}) \rightarrow \mathbf{X}_{\text{post}} \quad (1)$$

The function  $f$  is parameterized as a transformer that processes observational states and intervention specifications to predict post-intervention outcomes. Figure 1 illustrates how our implicit approach differs from explicit structure




Figure 1. Comparison of structure learning (left) versus effect prediction (right). In implicit mode, we bypass the structure learning pathway entirely, learning intervention effects directly through transformer attention without recovering the causal graph.

learning.

### 1.3. Contributions

Our work makes the following contributions:

- Conceptual:** We demonstrate that implicit learning is viable for causal inference, achieving competitive accuracy without graph recovery.
- Empirical:** On 5,000 zero-shot test cases, we outperform both PC Algorithm (1.34 vs 1.37 MAE) and NOTEARS (1.34 vs 1.40 MAE) while achieving best RMSE (4.71) among learnable methods.
- Scaling:** We successfully scale to 50 variables—more than double the typical limit of graph-based methods—with a systematic curriculum learning approach.
- Analysis:** A 768-dimensional model experiment (43% more parameters) reveals that the capacity limit stems from  $O(N^2)$  attention complexity rather than parameter count, informing future architectural directions.
- Practical:** Our method enables real-time deployment with 3.9ms inference time, 92 $\times$  faster than competitive ML baselines.

### 1.4. Results Overview

Our implicit approach achieves strong performance: (1) matches PC Algorithm accuracy within 1%, (2) outperforms NOTEARS by 4.3%, (3) achieves best RMSE/R<sup>2</sup> among all learnable methods, and (4) maintains only a 12% gap from Oracle performance (which has access to the true causal graph). Most critically, at 30-40 variables where baselines degrade significantly, our method maintains or improves accuracy, demonstrating superior scaling properties.

continue here

## 2. Related Work

### 2.1. Structure Learning Methods

**Constraint-based approaches.** The PC Algorithm (Spirtes et al., 2000) uses conditional independence testing to identify graph structure. While principled, it struggles with limited sample sizes and shows significant performance degradation beyond 20 variables. Our method achieves comparable accuracy (1.34 vs 1.37 MAE) without requiring independence tests or graph recovery.

**Score-based methods.** GES (Chickering, 2002) performs greedy search over graph space. These methods face scalability challenges due to the combinatorial search space.

**Continuous optimization.** NOTEARS (Zheng et al., 2018) reformulates DAG learning as continuous optimization with a novel acyclicity constraint. While innovative, our experiments show significant degradation at 40 variables (1.66 MAE) compared to our method (1.42 MAE), and the approach still requires explicit graph recovery.

### 2.2. Deep Learning for Causal Structure

Recent work applies neural networks to causal discovery: DAG-GNN (Yu et al., 2019), GraN-DAG (Lachapelle et al., 2020), and ENCO (Lippe et al., 2022). These methods still focus on learning graph structure first. Our approach fundamentally differs by learning intervention effects *without* explicit graph recovery, circumventing structural identifiability issues.

### 2.3. Treatment Effect Estimation

Methods for Individual Treatment Effect (ITE) and Conditional Average Treatment Effect (CATE) estimation (Shalit et al., 2017; Künzel et al., 2019) predict intervention outcomes. However, these typically predict effects on a single outcome variable given a treatment. Our method predicts effects on *all* system variables simultaneously, capturing propagated causal effects throughout the entire system.

### 2.4. Transformers for Structured Data

**Prior-Data Fitted Networks.** Müller et al. (Müller et al., 2021) demonstrated that transformers pre-trained on synthetic datasets can perform in-context Bayesian inference. TabPFN (Hollmann et al., 2022) extended this to tabular classification, achieving competitive performance through meta-learning on diverse synthetic tasks.

**Our approach.** While TabPFN focuses on supervised classification with pre-training, we adapt their key architectural innovations for causal prediction: (1) the MLP-based value embedding for continuous features, and (2) interleaved [Feature, Value] token representation that separates vari-




Figure 2. System overview showing the complete pipeline from data generation through training to evaluation. Data generation creates random DAGs with interventional samples. The 512-dim transformer with 12 blocks (8 attention heads, 2048 FFN dimension) learns via curriculum training (20–50 variables) with cosine annealing and LR reset. Evaluation uses zero-shot testing on unseen graphs.

able identity from observed values. However, we train from scratch on causal intervention data rather than using meta-learning, as our task requires understanding causal mechanisms that vary across different structural causal models rather than learning a universal prediction strategy.

### 3. Method

#### 3.1. Problem Formulation

We consider a system of  $N$  variables  $\mathbf{X} = (X_1, \dots, X_N) \in \mathbb{R}^N$  governed by a Structural Causal Model (SCM):

$$X_i = f_i(\text{PA}_i, \epsilon_i), \quad i = 1, \dots, N \quad (2)$$

where  $\text{PA}_i$  denotes the parents of  $X_i$  in the causal DAG  $\mathcal{G}$  and  $\epsilon_i$  represents exogenous noise.

**Intervention operator.** An intervention  $\text{do}(X_i = v)$  replaces the mechanism for  $X_i$  with a constant, yielding post-intervention distribution:

$$P(\mathbf{X} | \text{do}(X_i = v)) = \prod_{j \neq i} P(X_j | \text{PA}_j) \cdot \delta(X_i - v) \quad (3)$$

**Learning objective.** Given observational state  $\mathbf{X}_{\text{obs}}$  and intervention specification  $\mathcal{I} = \{\text{do}(X_{i_1} = v_1), \dots, \text{do}(X_{i_k} = v_k)\}$ , our goal is to learn:

$$f_\theta : (\mathbf{X}_{\text{obs}}, \mathcal{I}) \rightarrow \Delta \mathbf{X} \quad (4)$$

where  $\Delta \mathbf{X} = \mathbf{X}_{\text{post}} - \mathbf{X}_{\text{obs}}$  represents the causal effects, and the DAG structure  $\mathcal{G}$  is never explicitly recovered.

#### 3.2. Model Architecture

Our architecture consists of three components: embedding layer, transformer encoder, and prediction head.

**TabPFN-style value embedding.** Following TabPFN (Hollmann et al., 2022), we embed scalar values through a small MLP with  $2\times$  expansion rather




Figure 3. Model architecture showing the complete flow from input through embedding to output. Inputs (observational data and interventions) are embedded using TabPFN-style MLP with positional and type encodings, creating  $2N$  interleaved tokens. These flow through 12 transformer blocks (8-head attention, 2048 FFN dimension) with residual connections and layer normalization. The effect prediction head outputs intervention effects  $\Delta X$  for all variables.

than simple linear projection. This provides richer representations for continuous intervention values:

$$\mathbf{h}_1 = \text{GELU}(W_1 x_i) \quad W_1 \in \mathbb{R}^{2d \times 1} \quad (5)$$

$$\mathbf{e}_{\text{val}}(x_i) = \text{LayerNorm}(W_2 \mathbf{h}_1) \quad W_2 \in \mathbb{R}^{d \times 2d} \quad (6)$$

This design handles the diversity of intervention magnitudes better than direct embedding, improving MAE by 5% in ablations (Table 4).

**Complete embedding.** We create token representations through three components:

- **Variable ID:**  $\mathbf{e}_{\text{id}}^{(i)} \in \mathbb{R}^d$  (learned positional encoding)
- **Value:**  $\mathbf{e}_{\text{val}}(x_i)$  as defined above
- **Type:**  $\mathbf{e}_{\text{type}}^{(t)} \in \mathbb{R}^d$  where  $t \in \{0, 1\}$  indicates observed vs intervened

**Interleaved token strategy.** Rather than single token per variable, we create two tokens—a feature token encoding variable identity and a value token encoding the measurement—yielding sequence  $[\mathbf{e}_{\text{id}}^{(0)}, \mathbf{e}_{\text{val}}(x_0) + \mathbf{e}_{\text{type}}, \mathbf{e}_{\text{id}}^{(1)}, \mathbf{e}_{\text{val}}(x_1) + \mathbf{e}_{\text{type}}, \dots]$  with  $2N$  tokens total. This separation—inspired by TabPFN’s design for tabular data—allows the transformer to attend to “which variable” independently from “what value”, critical for learning intervention patterns. Ablations show 8% MAE improvement over single-token design. Figure 3 illustrates the architecture.

**Transformer encoder.** We use 12 transformer layers with pre-normalization:

$$\mathbf{h}' = \mathbf{h} + \text{MultiHeadAttn}(\text{LayerNorm}(\mathbf{h})) \quad (7)$$

$$\mathbf{h}_{\text{new}} = \mathbf{h}' + \text{FFN}(\text{LayerNorm}(\mathbf{h}')) \quad (8)$$

where FFN includes GELU activation and dropout. Pre-norm provides better gradient flow for deep networks. All

165 12 layers use 8-head attention with Flash Attention (Dao  
166 et al., 2022) optimization for efficiency.  
167

168 **Prediction head.** A linear layer maps value token embed-  
169 dings to effect predictions:

$$\Delta X_i = W_{\text{out}} \mathbf{h}_{2i+1}^{(12)} \quad (9)$$

170 where  $\mathbf{h}_{2i+1}^{(12)}$  extracts the value token for variable  $i$  from the  
171 final layer.  
172

### 3.3. Curriculum Learning with LR Reset

173 **Curriculum strategy.** We progressively train from simple  
174 to complex systems:  
175

$$L_0(20 \text{ vars}) \rightarrow L_1(21 \text{ vars}) \rightarrow \dots \rightarrow L_{30}(50 \text{ vars}) \quad (10)$$

182 At each level  $\ell$ , we train until validation MAE drops below  
183 threshold  $\theta_\ell$  (e.g., 0.15 for 20 vars, 0.45 for 50 vars), then  
184 save a checkpoint and advance.  
185

186 **Critical innovation: LR reset.** Upon ad-  
187 vancing from  $\ell$  to  $\ell + 1$ , we *reset* the opti-  
188 mizer:

189 Load weights from  $L_\ell$  checkpoint

190 **Reset:** optimizer  $\leftarrow$  AdamW(params, lr=1e-4)

191 Train on  $L_{\ell+1}$  (more variables)

192 This reset provides fresh optimizer momentum for in-  
193 creased complexity, enabling scaling from 35 to 50 vari-  
194 ables—our ablation studies show this is the single most  
195 critical component.  
196

197 **Loss function.** We minimize MSE on intervention effects:

$$\mathcal{L} = \frac{1}{NB} \sum_{b=1}^B \sum_{i=1}^N (\Delta X_i^{(b)} - \hat{\Delta X}_i^{(b)})^2 \quad (11)$$

202 with gradient clipping at norm 1.0 (critical for stability at  
203 40+ variables). Figure 4 illustrates the complete training  
204 workflow with curriculum advancement.  
205

### 3.4. Data Generation

206 We generate synthetic SCMs with known ground truth for  
207 evaluation:  
208

- 211 • **DAG:** Erdős-Rényi graph with topological ordering,  
212 density 0.15-0.30
- 213 • **Parameters:**  $w_{ij} \sim \text{Uniform}(-2, 2)$
- 214 • **Noise:**  $\epsilon_i \sim \mathcal{N}(0, 1)$
- 215 • **Interventions:** Random 1-3 variables per sample, val-  
216 ues  $\sim \mathcal{N}(0, 3)$




Figure 4. Training pipeline with curriculum learning. Starting at 20 variables, each epoch generates batches, performs forward/backward passes with gradient clipping (max\_norm=1.0), and validates. When MAE drops below threshold (and epochs  $\geq 5$ ), we save a checkpoint, advance to  $N + 1$  variables, and reset the learning rate to 1e-4 for fresh momentum. This continues until reaching 50 variables, enabling successful scaling beyond traditional methods.

Training uses 1,000 samples per variable count with diverse graph structures. While synthetic, this allows rigorous zero-shot evaluation and Oracle baseline comparison. Figure 5 illustrates the complete data generation pipeline.




Figure 5. Data generation pipeline. For each training sample, we: (1) generate a random DAG with specified variable count and edge density, (2) sample linear SCM parameters and noise distributions, (3) generate observational data from the SCM, (4) randomly select 1–3 variables for intervention with values drawn from  $\mathcal{N}(0, 9)$ , (5) compute post-intervention states by propagating through the SCM, and (6) create training batches containing observational states, intervention specifications, and post-intervention outcomes.

## 220 4. Experiments

### 221 4.1. Experimental Setup

222 **Test configurations.** We evaluate on 5,000 interventional  
 223 predictions across completely unseen graphs:

Config	Vars	Density	Graphs	Total
1	20	0.20	100	1,000
2	20	0.30	100	1,000
3	30	0.20	100	1,000
4	30	0.30	100	1,000
5	40	0.25	100	1,000

232 *Table 1.* Test configurations for zero-shot evaluation.

### 233 234 Baselines.

- **Oracle:** Uses true graph (upper bound)
- **PC Algorithm** (Spirtes et al., 2000): Constraint-based discovery
- **NOTEARS** (Zheng et al., 2018): Continuous optimization
- **Random Forest:** ML baseline
- **Linear Regression:** Simple baseline
- **MLP:** Neural network baseline

250 **Metrics.** MAE (primary), RMSE (prediction quality), R<sup>2</sup>  
 251 (model fit), inference time.

### 252 4.2. Main Results

255 Table 2 shows our comprehensive evaluation. Our method  
 256 achieves 1.34 MAE, matching PC Algorithm (1.37) within  
 257 1% and outperforming NOTEARS (1.40) by 4.3%.


259 Crucially, we achieve **best RMSE (4.71)** and **best R<sup>2</sup> (-0.02)**  
 260 among all learnable methods, indicating superior  
 261 prediction quality. The 12% gap from Oracle (1.34 vs  
 262 1.21 MAE) represents the cost of not having perfect graph  
 263 knowledge—a remarkably small price for avoiding expen-  
 264 sive graph search. Figure 6 illustrates the R<sup>2</sup> comparison  
 265 across methods.

### 266 4.3. Scaling Performance

268 Figure 8 reveals our key advantage: while baselines de-  
 269 grade at 30-40 variables, our method maintains strong per-  
 270 formance. At 30 variables, we achieve 1.32 MAE vs 1.41  
 271 for the best baseline—a clear lead. At 40 variables, this gap  
 272 widens: 1.42 vs 1.66 (NOTEARS), demonstrating success-  
 273 ful scaling where traditional methods fail.

Method	MAE ↓	RMSE ↓	R <sup>2</sup> ↑	Time
Oracle	<b>1.21</b>	<b>4.64</b>	<b>0.02</b>	5.2ms
<b>Ours</b>	<b>1.34</b>	<b>4.71</b>	<b>-0.02</b>	<b>3.9ms</b>
PC Algorithm	1.37	4.91	-0.12	6.1ms
NOTEARS	1.40	5.16	-0.28	7.3ms
Random Forest	1.34	4.95	-0.13	358ms
Linear Reg	1.36	4.91	-0.10	0.4ms
MLP	1.46	5.03	-0.16	0.3ms

232 *Table 2.* Overall performance on 5,000 zero-shot test cases. Bold  
 233 indicates best among learnable methods.



234 *Figure 6.* R<sup>2</sup> comparison across methods. Our approach achieves  
 235 the best R<sup>2</sup> among learnable methods, closest to Oracle perfor-  
 236 mance.

### 237 4.4. Ablation Studies

238 Table 4 analyzes key components. Removing LR reset  
 239 causes failure beyond 35 variables (most critical). Gradi-  
 240 ent clipping prevents training instability. Interleaved tokens  
 241 improve MAE by 8%, while TabPFN-style embedding con-  
 242 tributes 5% improvement.

### 243 4.5. Scaling Analysis: 768-Dimensional Model

244 To investigate capacity limits, we trained a 768-  
 245 dimensional variant (50% wider, 86M parameters vs 60M).  
 246 Despite 43% more parameters, capacity improved only  
 247 from 46 to 51 variables (+11%). Both models encountered  
 248 similar limits around 50-55 variables.

249 Parameter efficiency actually *decreased*: 512-dim achieved  
 250 59% of theoretical capacity (46/78) while 768-dim  
 251 achieved only 53% (51/96). This demonstrates the bot-  
 252 tleneck is not parameter count but rather  $O(N^2)$  full atten-  
 253 tion complexity on  $2N$  interleaved tokens. At 46 vari-  
 254 ables:  $92^2 = 8,464$  attention operations; at 52 variables:  
 255  $104^2 = 10,816$  operations—a 28% increase that causes  
 256 failure in both models.

257 This negative result is scientifically valuable: it rules out  
 258 simple parameter scaling and motivates sparse attention  
 259 mechanisms as a more promising direction.




Figure 7. Inference time comparison. Our method achieves real-time performance (3.9ms) significantly faster than Random Forest while maintaining competitive accuracy.




Figure 8. MAE vs number of variables. Our method (blue) maintains performance while NOTEARS (orange) degrades significantly at 40 variables.

## 5. Discussion

### 5.1. When to Use Implicit vs Explicit Discovery

### 5.2. Why Implicit Learning Works

Our results challenge the intuition that graph recovery is necessary for intervention prediction. We hypothesize that transformer attention implicitly learns *functional* rather than *structural* representations. While explicit methods learn  $\mathcal{G}$  then compute  $\mathbb{E}[X_j \mid \text{do}(X_i)]$  via graph operations, our model learns a direct mapping from intervention specifications to effects.

This works because: (1) the training data (1,000 samples  $\times$  diverse graphs  $\times$  curriculum levels) provides rich coverage of causal patterns, allowing the model to learn the invariant structure of do operators across different graphs, and (2) TabPFN-style interleaved tokens enable the model to represent both “which variables interact” (via attention patterns) and “how strong is the effect” (via value representations) without committing to a discrete adjacency matrix.

The 12% gap from Oracle (1.34 vs 1.21 MAE) likely represents information loss from not having perfect graph knowledge, but this cost is acceptable given the computational savings and scaling advantages.

Vars	Ours	Oracle	NOTEARS	PC
20	1.25	1.12	1.24	1.27
30	<b>1.32</b>	1.23	1.44	1.41
40	<b>1.42</b>	1.27	1.66	1.48
46	<b>1.45</b>	1.30	-	-
50	<b>1.48</b>	1.32	-	-

Table 3. MAE by variable count. Only our method scales beyond 40 variables.




Figure 9. MAE heatmap across different variable counts and graph densities. Our method shows consistent performance across configurations.

### 5.3. Limitations and Future Work

#### Current limitations:

- **Synthetic data only.** While synthetic SCMs provide rigorous evaluation with known ground truth, real-world validation is critical. We expect strong transfer given that our model learns functional relationships rather than dataset-specific patterns, but empirical confirmation is needed.
- **Linear SCMs.** Real-world causal systems often involve nonlinearities. Extension to nonlinear mechanisms (neural SCMs, additive noise models) is straightforward but requires validation that TabPFN-style embeddings can handle complex functional forms.
- **Capacity limit at 50 variables.** The  $O(N^2)$  attention bottleneck limits applicability to medium-scale systems. Sparse attention mechanisms offer a clear path forward (see below).
- **Black box nature.** Unlike explicit methods that produce interpretable graphs, our model’s learned representations are opaque. Attention visualization and probing studies could provide partial interpretability.

Configuration	MAE	Max Vars
Full model	<b>1.34</b>	<b>50</b>
w/o LR reset	1.52	35
w/o gradient clip	unstable	-
w/o interleaved tokens	1.45	46
w/o TabPFN embedding	1.41	46

Table 4. Ablation study on key components.




Figure 10. Error distributions across different system sizes. Our method maintains tight error distributions even as complexity scales.

- **Training cost.** Full curriculum training (20→50 vars)

Scenario	Recommendation
Prediction, 30+ vars	<b>Implicit:</b> Better scaling; maintains accuracy
Prediction, <20 vars	<b>Either:</b> Both competitive
Need mechanism	<b>Explicit:</b> Interpretable graph
Real-time deploy	<b>Implicit:</b> Fast inference (<4ms)
Limited data (~100)	<b>Explicit:</b> More sample efficient
Rich data (~1000+)	<b>Implicit:</b> Amortizes learning
Safety-critical	<b>Explicit:</b> Aids verification
High-throughput	<b>Implicit:</b> Constant-time prediction

Table 5. Decision guide for method selection. Implicit methods excel at prediction-focused tasks with sufficient data and scale, while explicit methods remain preferable when interpretability or sample efficiency is critical.

requires  $\sim 8$  hours on  $2 \times$ A100 GPUs. However, this is a one-time cost; inference remains fast (3.9ms).

**Scaling beyond 50 variables.** Our 768-dim experiment (Section 4.5) revealed that capacity limits stem from  $O(N^2)$  attention complexity rather than parameter count. Increasing model width gave only marginal improvements (+11% capacity for +43% parameters), confirming that the bottleneck is architectural, not representational.

Concrete path forward:

- **Sparse attention** (Child et al., 2019): Local or block-sparse patterns reduce complexity to  $O(N\sqrt{N})$  or  $O(N \log N)$ . We expect this enables 80–100 variables on current hardware with minimal accuracy loss, as causal effects are often localized.
- **Hierarchical models:** Process variables in groups (e.g., organ systems in medical data), then combine group-level predictions. Expected capacity: 100–200 variables.
- **Linear attention** (Katharopoulos et al., 2020): Reduces complexity to  $O(N)$  but may sacrifice modeling power. Worth investigating as a simple baseline.

All three approaches address the  $O(N^2)$  bottleneck more fundamentally than naive parameter scaling.

**Real-world validation.** Promising domains include:

- **Healthcare:** Predicting patient outcomes under treatment interventions from electronic health records (20–40 biomarkers). Our fast inference enables real-time clinical decision support.
- **Economics:** Estimating policy intervention effects in macroeconomic systems (GDP, inflation, employment, etc.). The ability to handle 30+ economic indicators is critical.

- **Climate modeling:** Predicting climate variable responses to mitigation strategies. Scaling to 50+ variables enables whole-Earth system modeling.
- **Robotics:** Model-based reinforcement learning where intervention predictions guide action selection. Real-time inference (<4ms) crucial for control loops.

**Practical deployment.** Beyond accuracy, deployment requires: (1) uncertainty quantification (can be added via ensembles or Bayesian neural networks), (2) handling missing values (straightforward via type embeddings), and (3) continual learning as new data arrives (curriculum learning transfers naturally to online settings).

## 6. Conclusion

We introduced implicit causal discovery as an alternative to traditional graph-first methods, demonstrating that intervention effects can be learned directly via transformer attention mechanisms. Our approach achieves competitive accuracy with explicit causal discovery methods (matching PC Algorithm, outperforming NOTEARS) while scaling to 50 variables—more than double the typical limit of graph-based methods. With best RMSE/R<sup>2</sup> among learnable methods and only a 12% gap from Oracle, we show that graph recovery is not prerequisite to accurate intervention prediction.

This work challenges the assumption that explicit structure learning is necessary for causal inference, opening new avenues for practical intervention prediction in high-dimensional systems. As evidenced by our scaling analysis, the path forward involves architectural innovations (sparse attention, hierarchical modeling) rather than simple parameter expansion. Such developments may enable implicit causal discovery in domains with hundreds of variables, previously beyond reach of traditional methods.

## References

- Chickering, D. M. Learning bayesian networks is np-complete. *Learning from data*, pp. 121–130, 1996.
- Chickering, D. M. Optimal structure identification with greedy search. *Journal of machine learning research*, 3 (Nov):507–554, 2002.
- Child, R., Gray, S., Radford, A., and Sutskever, I. Generating long sequences with sparse transformers. *arXiv preprint arXiv:1904.10509*, 2019.
- Dao, T., Fu, D. Y., Ermon, S., Rudra, A., and Ré, C. FlashAttention: Fast and memory-efficient exact attention with io-awareness. In *Advances in Neural Information Processing Systems*, 2022.
- Hollmann, N., Müller, S., Eggensperger, K., and Hutter, F. TabPFN: A transformer that solves small tabular classification problems in a second. In *International Conference on Learning Representations*, 2022.
- Katharopoulos, A., Vyas, A., Pappas, N., and Fleuret, F. Transformers are rnns: Fast autoregressive transformers with linear attention. In *International conference on machine learning*, pp. 5156–5165. PMLR, 2020.
- Künzel, S. R., Sekhon, J. S., Bickel, P. J., and Yu, B. Metalearners for estimating heterogeneous treatment effects using machine learning. In *Proceedings of the national academy of sciences*, volume 116, pp. 4156–4165, 2019.
- Lachapelle, S., Brouillard, P., Deleu, T., and Lacoste-Julien, S. Gradient-based neural dag learning. In *International Conference on Learning Representations*, 2020.
- Lippe, P., Cohen, T., and Gavves, E. Efficient neural causal discovery without acyclicity constraints. In *International Conference on Learning Representations*, 2022.
- Müller, S., Hollmann, N., Arango, S. P., Grabocka, J., and Hutter, F. Transformers can do bayesian inference. *arXiv preprint arXiv:2112.10510*, 2021.
- Pearl, J. *Causality*. Cambridge university press, 2009.
- Shalit, U., Johansson, F. D., and Sontag, D. Estimating individual treatment effect: generalization bounds and algorithms. In *International Conference on Machine Learning*, pp. 3076–3085. PMLR, 2017.
- Spirites, P., Glymour, C. N., and Scheines, R. *Causation, prediction, and search*. MIT press, 2000.
- Yu, Y., Chen, J., Gao, T., and Yu, M. Dag-gnn: Dag structure learning with graph neural networks. In *International Conference on Machine Learning*, pp. 7154–7163. PMLR, 2019.
- Zheng, X., Aragam, B., Ravikumar, P. K., and Xing, E. P. Dags with no tears: Continuous optimization for structure learning. In *Advances in Neural Information Processing Systems*, pp. 9472–9483, 2018.

## A. Appendix

### A.1. Architecture Details

#### A.1.1. MODEL CONFIGURATION

Our transformer architecture consists of 12 encoder layers with the following specifications:

Component	Specification	Parameters
<b>Embedding Dimension</b>	$d_{model} = 512$	-
<b>Transformer Layers</b>	12 layers	-
<b>Attention Heads</b>	8 heads	-
<b>FFN Hidden Dimension</b>	$d_{ff} = 2048$	-
<b>Dropout</b>	0.1	-
<b>Activation</b>	GELU	-
<b>Normalization</b>	Pre-norm (LayerNorm)	-
<b>Total Parameters</b>	-	$\sim 60M$

Table 6. Core transformer architecture configuration.

#### A.1.2. LAYER-BY-LAYER BREAKDOWN

Component	Parameters	% Total
<b>Embedding Layer</b>		
Variable ID Embedding	26K	0.04%
Value Embedding MLP	0.5M	0.8%
Type Embedding	1K	0.002%
<i>Subtotal</i>	<i>8M</i>	<i>13%</i>
<b>Transformer Layers (11x)</b>		
Attention per layer	1.05M	-
FFN per layer	2.10M	-
<i>Subtotal</i>	<i>35M</i>	<i>58%</i>
<b>Gumbel Layer (12th)</b>	3.2M	5%
<b>Output Head</b>	512	0.001%
<b>LayerNorms</b>	12K	0.02%
<b>Total</b>	<b><math>\sim 60M</math></b>	<b>100%</b>

Table 7. Parameter distribution across model components.

#### A.1.3. EMBEDDING STRATEGY

We employ a TabPFN-style embedding with interleaved [Feature, Value] tokens:

- **Variable ID:** Learned positional encoding for each variable
- **Value Embedding:** 2-layer MLP ( $1 \rightarrow 1024 \rightarrow 512$ ) with GELU and LayerNorm
- **Type Embedding:** Binary indicator (observed=0, intervened=1)
- **Token Sequence:**  $[F_0, V_0, F_1, V_1, \dots, F_{N-1}, V_{N-1}]$  (2N tokens total)

## A.2. Training Configuration

### A.2.1. HYPERPARAMETERS

Category	Parameter	Value
<b>Optimizer</b>	AdamW	-
	Learning Rate	1e-4
	Weight Decay	0.01
	$\beta_1, \beta_2$	0.9, 0.999
	$\epsilon$	1e-8
<b>Batch Configuration</b>	Per-GPU Batch Size	2
	Number of GPUs	4
	Gradient Accumulation	8
	Effective Batch Size	64
<b>Regularization</b>	Gradient Clipping	1.0
	Dropout Rate	0.1
<b>LR Schedule</b>	Scheduler	CosineAnnealingWarmRestarts
	$T_0$ (restart every)	50 epochs
	$\eta_{min}$	1e-6

Table 8. Training hyperparameters and optimization settings.

### A.2.2. CURRICULUM LEARNING

Variables	MAE Threshold	Typical Epochs	Memory (Training)
20-25	< 0.15	~30	~8 GB
26-30	< 0.25	~40	~16 GB
31-35	< 0.30	~50	~24 GB
36-40	< 0.35	~50	~24 GB
41-50	< 0.45	~60-80	~48 GB

Table 9. Curriculum learning configuration with MAE thresholds and convergence patterns.

**Critical Strategy:** Learning rate reset to 1e-4 when advancing curriculum levels. This proved essential for scaling beyond 35 variables.

## A.3. Performance Metrics

Vars	Method	MAE	RMSE	R <sup>2</sup>
20	Ours	1.36	4.11	-0.046
	PC Algorithm	1.11	3.94	0.038
	Oracle	1.12	3.95	0.033
30	Ours	1.23	4.12	0.008
	PC Algorithm	1.21	4.15	-0.006
	Oracle	1.15	4.10	0.016
40	Ours	1.42	4.65	-0.028
	PC Algorithm	1.36	4.68	-0.042
	Oracle	1.22	4.57	0.004

Table 10. Performance comparison by variable count on zero-shot test cases.

495 **A.4. Additional Results**496 **A.4.1. MEMORY REQUIREMENTS**


---

Variables	Model Weights	Activations	Total (Train)	Inference
20	240 MB	2 GB	~3 GB	~2 GB
30	240 MB	4 GB	~5 GB	~4 GB
40	240 MB	8 GB	~9 GB	~6 GB
50	240 MB	24 GB	~48 GB	~9 GB

504 *Table 11.* Memory requirements by system size. Model weights  
 505 remain constant; activation memory scales with  $O(N^2)$  attention  
 506 complexity.

507 **A.4.2. INFERENCE SPEED**

508 Our method achieves real-time inference with mean time of  
 509 3.9ms ( $\pm 2.1$ ms std) averaged across 5,000 test cases. This  
 510 is 92 $\times$  faster than Random Forest (358ms) and competi-  
 511 tive with traditional causal discovery methods (PC: 6.1ms,  
 512 NOTEARS: 7.3ms).

513 **A.5. Reproducibility**

514 **Hardware:** 4 $\times$  NVIDIA A100 40GB GPUs

515 **Software:** Python 3.10, PyTorch 2.1.0, CUDA 11.8

516 **Training Command:**

```
517 torchrun --nproc_per_node=4 train.py \
518   --mode implicit --epochs 500 \
519   --batch_size 2 --lr 1e-4 \
520   --min_vars 20 --max_vars 50
```

521 **Code:** Available at <https://github.com/MeiisamMahmoodii/ISD-CP-Final>